

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
df = pd.read_csv('insurance.csv')
df.head()
```

Out[2]:

| | age | sex | bmi | children | smoker | region | charges | insuranceclaim |
|---|-----|-----|--------|----------|--------|--------|-------------|----------------|
| 0 | 19 | 0 | 27.900 | 0 | 1 | 3 | 16884.92400 | 1 |
| 1 | 18 | 1 | 33.770 | 1 | 0 | 2 | 1725.55230 | 1 |
| 2 | 28 | 1 | 33.000 | 3 | 0 | 2 | 4449.46200 | 0 |
| 3 | 33 | 1 | 22.705 | 0 | 0 | 1 | 21984.47061 | 0 |
| 4 | 32 | 1 | 28.880 | 0 | 0 | 1 | 3866.85520 | 1 |

In [3]:

```
# checking first five rows
df.head()
```

Out[3]:

| | age | sex | bmi | children | smoker | region | charges | insuranceclaim |
|---|-----|-----|--------|----------|--------|--------|-------------|----------------|
| 0 | 19 | 0 | 27.900 | 0 | 1 | 3 | 16884.92400 | 1 |
| 1 | 18 | 1 | 33.770 | 1 | 0 | 2 | 1725.55230 | 1 |
| 2 | 28 | 1 | 33.000 | 3 | 0 | 2 | 4449.46200 | 0 |
| 3 | 33 | 1 | 22.705 | 0 | 0 | 1 | 21984.47061 | 0 |
| 4 | 32 | 1 | 28.880 | 0 | 0 | 1 | 3866.85520 | 1 |

In [4]:

```
# checking last five rows
df.tail()
```

Out[4]:

| | age | sex | bmi | children | smoker | region | charges | insuranceclaim |
|------|-----|-----|-------|----------|--------|--------|------------|----------------|
| 1333 | 50 | 1 | 30.97 | 3 | 0 | 1 | 10600.5483 | 0 |
| 1334 | 18 | 0 | 31.92 | 0 | 0 | 0 | 2205.9808 | 1 |
| 1335 | 18 | 0 | 36.85 | 0 | 0 | 2 | 1629.8335 | 1 |
| 1336 | 21 | 0 | 25.80 | 0 | 0 | 3 | 2007.9450 | 0 |
| 1337 | 61 | 0 | 29.07 | 0 | 1 | 1 | 29141.3603 | 1 |

In [5]:

```
# shape of dataset

print('number of rows :',df.shape[0])
print('number of columns :',df.shape[1])
```

```
number of rows : 1338
number of columns : 8
```

In [6]:

```
# total number of datapoints in our dataset

print('size of our dataset :',df.size)
```

```
size of our dataset : 10704
```

In [7]:

```
df.isnull().sum()
```

Out[7]:

```
age                0
sex                0
bmi                0
children           0
smoker             0
region             0
charges            0
insuranceclaim     0
dtype: int64
```

In [8]:

```
# checking for information of our dataset
# number of columns
# type of columns
# their datatypes

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   1338 non-null   int64
 1   sex                   1338 non-null   int64
 2   bmi                   1338 non-null   float64
 3   children              1338 non-null   int64
 4   smoker                1338 non-null   int64
 5   region                1338 non-null   int64
 6   charges                1338 non-null   float64
 7   insuranceclaim        1338 non-null   int64
dtypes: float64(2), int64(6)
memory usage: 83.8 KB
```

Insight :

1. there are total 8 columns

2. there are no null values
3. 6 columns are of int64 datatype
4. 2 columns are of float64 datatype
5. there are no columns with object or boolean datatype

In [9]:

```
df.describe()
```

Out[9]:

| | age | sex | bmi | children | smoker | region | charges |
|--------------|-------------|-------------|-------------|-------------|-------------|-------------|--------------|
| count | 1338.000000 | 1338.000000 | 1338.000000 | 1338.000000 | 1338.000000 | 1338.000000 | 1338.000000 |
| mean | 39.207025 | 0.505232 | 30.663397 | 1.094918 | 0.204783 | 1.515695 | 13270.422265 |
| std | 14.049960 | 0.500160 | 6.098187 | 1.205493 | 0.403694 | 1.104885 | 12110.011237 |
| min | 18.000000 | 0.000000 | 15.960000 | 0.000000 | 0.000000 | 0.000000 | 1121.873900 |
| 25% | 27.000000 | 0.000000 | 26.296250 | 0.000000 | 0.000000 | 1.000000 | 4740.287150 |
| 50% | 39.000000 | 1.000000 | 30.400000 | 1.000000 | 0.000000 | 2.000000 | 9382.033000 |
| 75% | 51.000000 | 1.000000 | 34.693750 | 2.000000 | 0.000000 | 2.000000 | 16639.912515 |
| max | 64.000000 | 1.000000 | 53.130000 | 5.000000 | 1.000000 | 3.000000 | 63770.428010 |

insights :

1. min age is 18 and max age is 64
2. maximum customers are of mid age between 35-50
3. min claim settlement is 1121
4. max claim settlement is 63770
5. almost equal number of male and female customers
6. if the customer has higher bmi and is smoker then there is high probability of claim
7. many of the customers has 1-2 children

In []:

creating a base model without any changes

Splitting the dataset between independent variables and dependent variable

In [10]:

```
X_1 = df.drop(['insuranceclaim'],axis = 1)
y_1 = df['insuranceclaim']
```

In [11]:

```
# splitting data into train and test

from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test = train_test_split(X_1,y_1,test_size = 0.20,random_state = 42)
```

In [12]:

```
from sklearn.linear_model import LogisticRegression
from sklearn import svm
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score
```

In [13]:

```
# Ogistic regression

lr1 = LogisticRegression()

# fit on data

lr1.fit(X_train,y_train)
```

Out[13]:

```
LogisticRegression()
```

In [14]:

```
# predict

ypr1 = lr1.predict(X_test)
```

In [15]:

```
# accuracy

ac1 = accuracy_score(y_test,ypr1)
ac1
```

Out[15]:

```
0.7947761194029851
```

svm

In [16]:

```
# instance

svm1 = svm.SVC()

# fit on data

svm1.fit(X_train,y_train)
```

Out[16]:

SVC()

In [17]:

```
# predict

ypr2 = svm1.predict(X_test)

# accuracy

ac2 = accuracy_score(y_test,ypr2)
ac2
```

Out[17]:

0.6082089552238806

random forest classifier

In [18]:

```
# crereating instances

rf1 = RandomForestClassifier()

# fit on data

rf1.fit(X_train,y_train)
```

Out[18]:

RandomForestClassifier()

In [19]:

```
# predict

ypr3 = rf1.predict(X_test)

# accuracy

ac3 = accuracy_score(y_test,ypr3)
ac3
```

Out[19]:

0.914179104477612

Decision Tree Classifier

In [20]:

```
# creating instance

dt1 = DecisionTreeClassifier()

# fit on data

dt1.fit(X_train,y_train)
```

Out[20]:

```
DecisionTreeClassifier()
```

In [21]:

```
# predict

ypr4 = dt1.predict(X_test)

# accuracy score

ac4 = accuracy_score(y_test,ypr4)
ac4
```

Out[21]:

```
0.9776119402985075
```

Gradient boosting

In [22]:

```
# creating instance

gbc1 = GradientBoostingClassifier()

# fit on data

gbc1.fit(X_train,y_train)
```

Out[22]:

```
GradientBoostingClassifier()
```

In [23]:

```
# predict

ypr5 = gbc1.predict(X_test)

# accuracy

ac5 = accuracy_score(y_test,ypr5)
ac5
```

Out[23]:

```
0.9701492537313433
```

K-Nearest Neighbor

In [24]:

```
# creating instance

knn1 = KNeighborsClassifier()

# fit on data

knn1.fit(X_train,y_train)
```

Out[24]:

```
KNeighborsClassifier()
```

In [25]:

```
# predict

ypr6 = knn1.predict(X_test)

# accuracy score

ac6 = accuracy_score(y_test,ypr6)
ac6
```

C:\Users\WOLF\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

Out[25]:

```
0.6380597014925373
```

In [26]:

```
f_df = pd.DataFrame({'models':['lr1','svm1','rf1','dt1','gbc1','knn1'],
                     'accuracy':[ac1,ac2,ac3,ac4,ac5,ac6]})
```

In [27]:

```
f_df
```

Out[27]:

| | models | accuracy |
|---|--------|----------|
| 0 | lr1 | 0.794776 |
| 1 | svm1 | 0.608209 |
| 2 | rf1 | 0.914179 |
| 3 | dt1 | 0.977612 |
| 4 | gbc1 | 0.970149 |
| 5 | knn1 | 0.638060 |

In [28]:

```
sns.barplot(f_df['models'],f_df['accuracy'])  
plt.title('base_models vs base_accuracy')
```

C:\Users\WOLF\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

Out[28]:

Text(0.5, 1.0, 'base_models vs base_accuracy')



Insight :

1. As we can see our model is showing very high accuracy
2. But as this is just a base model without any alterations using eda and others
3. therefore it won't be a great idea to use it. So, that we will make another one

In [29]:


```
df.columns
```

Out[29]:

```
Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges',  
      'insuranceclaim'],  
      dtype='object')
```


In [30]:

```
columns = ['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges', 'insuranceclaim']
for i in columns:
    plt.figure()
    sns.boxplot(df[i])
```



In [31]:

```
np.where(df['charges'] >= 35000)
```

Out[31]:

```
(array([ 14,  19,  23,  29,  30,  34,  38,  39,  49,  53,  55,
        82,  84,  86,  94, 109, 123, 146, 158, 161, 175, 185,
       203, 240, 242, 251, 252, 254, 256, 263, 265, 271, 281,
       288, 292, 298, 312, 322, 327, 328, 330, 338, 373, 377,
       381, 420, 421, 422, 441, 476, 488, 500, 524, 530, 543,
       549, 558, 569, 577, 587, 609, 615, 621, 629, 665, 667,
       668, 674, 677, 682, 697, 706, 725, 736, 738, 739, 742,
       759, 803, 819, 826, 828, 842, 845, 850, 852, 856, 860,
       883, 893, 901, 917, 947, 951, 953, 956, 958, 1012, 1021,
      1022, 1031, 1036, 1037, 1047, 1049, 1062, 1070, 1090, 1096, 1111,
      1117, 1118, 1122, 1124, 1139, 1146, 1152, 1156, 1186, 1206, 1207,
      1218, 1230, 1240, 1241, 1249, 1284, 1288, 1300, 1301, 1303, 1313,
      1323], dtype=int64),)
```

In [32]:

```
np.where(df['bmi'] >= 46)
```

Out[32]:

```
(array([ 116,  286,  401,  438,  454,  543,  547,  549,  660,  847,  860,
        930,  941, 1047, 1088, 1317], dtype=int64),)
```

Insight :

As we can see there are some outliers in these features but as they are very low we won't consider doing any operations with it. I am not going to drop the outliers of charges as we already have as they are very low in numbers.

In []:

New Model

In [33]:

```
X = df.drop(['insuranceclaim'],axis = 1)
y = df['insuranceclaim']
```

In [34]:

```
# check for equal distribution
```


```
sns.countplot(y)
```

C:\Users\WOLF\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[34]:

<AxesSubplot:xlabel='insuranceclaim', ylabel='count'>



In [35]:

```
# handling imbalanced dataset with smote
# smote synthesizes new minority instances (it creates artificial datapoints that are slightly
!pip install imblearn

from imblearn.over_sampling import SMOTE
```

Requirement already satisfied: imblearn in c:\users\wolf\anaconda3\lib\site-packages (0.0)

Requirement already satisfied: imbalanced-learn in c:\users\wolf\anaconda3\lib\site-packages (from imblearn) (0.10.1)

Requirement already satisfied: numpy>=1.17.3 in c:\users\wolf\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.21.5)

Requirement already satisfied: joblib>=1.1.1 in c:\users\wolf\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.2.0)

Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\wolf\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.0.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\wolf\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (2.2.0)

Requirement already satisfied: scipy>=1.3.2 in c:\users\wolf\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.9.1)

In [36]:

```
X_res,y_res = SMOTE().fit_resample(X,y)
```

In [37]:

```
y_res.value_counts()
```

Out[37]:

```
1    783
```

```
0    783
```

```
Name: insuranceclaim, dtype: int64
```

splitting between training and test data

In [38]:

```
# importing train_test_split for data splitting

from sklearn.model_selection import train_test_split
```

In [39]:

```
X_train,X_test,y_train,y_test = train_test_split(X_res,y_res,test_size = 0.20,random_state = 42)
```

In [40]:

```
print('X_train shape : ',X_train.shape)
print('X_test shape : ',X_test.shape)
print('y_train shape : ',y_train.shape)
print('y_test shape : ',y_test.shape)
```

```
X_train shape : (1252, 7)
```

```
X_test shape : (314, 7)
```

```
y_train shape : (1252,)
```

```
y_test shape : (314,)
```

In [41]:

```
# Feature scaling

from sklearn.preprocessing import StandardScaler

std = StandardScaler()

X_train = std.fit_transform(X_train)
X_test = std.transform(X_test)
```

In [42]:

```
# we have to check to confirm scaling is done

X_train
```

Out[42]:

```
array([[ -0.5639978 , -0.92902919, -0.18429644, ..., -0.45916416,
         1.43046451, -0.76494058],
       [  1.31553296,  1.07639244,  2.25251124, ..., -0.45916416,
         1.43046451, -0.08861016],
       [ -1.28689425,  1.07639244,  1.11088353, ..., -0.45916416,
        -0.44094778, -0.92452615],
       ...,
       [ -0.13025994, -0.92902919,  2.90343932, ...,  2.17787035,
         1.43046451,  2.90040653],
       [ -1.50376319,  1.07639244, -1.15130979, ..., -0.45916416,
        -1.37665393, -0.94290046],
       [  1.17095367,  1.07639244, -0.05077273, ..., -0.45916416,
         1.43046451, -0.20643984]])
```

As we can see that our target variable has two classes therefore this is a classification problem

Now we will import different classification algorithms for model prediction

In [43]:

```
from sklearn.linear_model import LogisticRegression
from sklearn import svm
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
```

Logistic Regression

In [44]:

```
# creating instance

lr = LogisticRegression()

# fit on data

lr.fit(X_train,y_train)
```

Out[44]:

```
LogisticRegression()
```

In [45]:

```
# prediction
y_pred1 = lr.predict(X_test)
```

In [46]:

```
# checking model prdiction accuracy

from sklearn.metrics import accuracy_score
```

In [47]:

```
acc1 = accuracy_score(y_test,y_pred1)
acc1
```

Out[47]:

```
0.8789808917197452
```

In [48]:

```
# as our data set is imbalanced it is very fdangeroyus to use sccuracy score
# therefore we will use precision,recall and f1 score

from sklearn.metrics import precision_score,recall_score,f1_score
```

In [49]:

```
# as we know precision is  $P = TP/(FP + TP)$ 

ps1 = precision_score(y_test,y_pred1)
ps1
```

Out[49]:

```
0.9060402684563759
```

In [50]:

```
# as we know recalLL is  $R = TP/(TP+FN)$ 

rs1 = recall_score(y_test,y_pred1)
rs1
```

Out[50]:

```
0.8490566037735849
```

In [51]:

```
# f1 score

f1_s1 = f1_score(y_test,y_pred1)
f1_s1
```

Out[51]:

0.8766233766233767

Support vector machine

In [52]:

```
# creating instance

svm = svm.SVC()

# fit on data

svm.fit(X_train,y_train)
```

Out[52]:

SVC()

In [53]:

```
# prediction

y_pred2 = svm.predict(X_test)
```

In [54]:

```
# checking accuracy

acc2 = accuracy_score(y_test,y_pred2)
acc2
```

Out[54]:

0.8949044585987261

In [55]:

```
# precision

ps2 = precision_score(y_test,y_pred2)
ps2
```

Out[55]:

0.9256756756756757

In [56]:

```
# recall

rs2 = recall_score(y_test,y_pred2)
rs2
```

Out[56]:

0.8616352201257862

In [57]:

```
# f1 score

f1_s2 = f1_score(y_test,y_pred2)
f1_s2
```

Out[57]:

0.8925081433224755

Random Forest Classifier

In [58]:

```
# creating instance

rfc = RandomForestClassifier()

# fit on data

rfc.fit(X_train,y_train)
```

Out[58]:

RandomForestClassifier()

In [59]:

```
# prediction

y_pred3 = rfc.predict(X_test)
```

In [60]:

```
# accuracy

acc3 = accuracy_score(y_test,y_pred3)
acc3
```

Out[60]:

0.9426751592356688

In [61]:

```
# precision
```

```
ps3 = precision_score(y_test,y_pred3)
ps3
```

Out[61]:

0.9795918367346939

In [62]:

```
# recall
```

```
rs3 = recall_score(y_test,y_pred3)
rs3
```

Out[62]:

0.9056603773584906

In [63]:

```
# f1 score
```

```
f1_s3 = f1_score(y_test,y_pred3)
f1_s3
```

Out[63]:

0.9411764705882353

Decision Tree Classifier

In [64]:

```
# creating instance
```

```
dtc = DecisionTreeClassifier()
```

```
# fit on data
```

```
dtc.fit(X_train,y_train)
```

Out[64]:

DecisionTreeClassifier()

In [65]:

```
# prediction
```

```
y_pred4 = dtc.predict(X_test)
```


In [66]:

```
# accuracy
```

```
acc4 = accuracy_score(y_test,y_pred4)  
acc4
```

Out[66]:

0.945859872611465

In [67]:

```
# precision
```

```
ps4 = precision_score(y_test,y_pred4)  
ps4
```

Out[67]:

0.9551282051282052

In [68]:

```
# recall
```

```
rs4 = recall_score(y_test,y_pred4)  
rs4
```

Out[68]:

0.9371069182389937

In [69]:

```
# f1 score
```

```
f1_s4 = f1_score(y_test,y_pred4)  
f1_s4
```

Out[69]:

0.946031746031746

K-Nearest Classifier

In [70]:

```
# creating instance
```

```
knn = KNeighborsClassifier()
```

```
# fit on data
```

```
knn.fit(X_train,y_train)
```

Out[70]:

KNeighborsClassifier()

In [71]:

predict

y_pred5 = knn.predict(X_test)

C:\Users\WOLF\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

In [72]:

accuracy

```
acc5 = accuracy_score(y_test, y_pred5)
acc5
```

Out[72]:

0.8949044585987261

As we have to assign values of k neighbors

In [73]:

score = []

for k in range(1,40):

```
knn = KNeighborsClassifier(n_neighbors = k)
```

```
knn.fit(X_train, y_train)
```

```
y_pred = knn.predict(X_test)
```

```
score.append(accuracy_score(y_test, y_pred))
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\WOLF\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\WOLF\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\WOLF\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along.

In [74]:

```
score
```

Out[74]:

```
[0.856687898089172,  
 0.8630573248407644,  
 0.8757961783439491,  
 0.8821656050955414,  
 0.8949044585987261,  
 0.8885350318471338,  
 0.8949044585987261,  
 0.8853503184713376,  
 0.8885350318471338,  
 0.8885350318471338,  
 0.8821656050955414,  
 0.8789808917197452,  
 0.8694267515923567,  
 0.8630573248407644,  
 0.8662420382165605,  
 0.8789808917197452,  
 0.8789808917197452,  
 0.8757961783439491,  
 0.8789808917197452,  
 0.8789808917197452,  
 0.8789808917197452,  
 0.8757961783439491,  
 0.8789808917197452,  
 0.8694267515923567,  
 0.8821656050955414,  
 0.8821656050955414,  
 0.8757961783439491,  
 0.8598726114649682,  
 0.8598726114649682,  
 0.8630573248407644,  
 0.8598726114649682,  
 0.8598726114649682,  
 0.8598726114649682,  
 0.8598726114649682,  
 0.8535031847133758,  
 0.8535031847133758,  
 0.8535031847133758,  
 0.8535031847133758,  
 0.856687898089172,  
 0.8535031847133758]
```

In [75]:

```
# as we are getting better accuracy for 5 neighbors as we can see in score
# therefore we will choose 5 neighbors

for k in range(1,40):
    knn = KNeighborsClassifier(n_neighbors = 5)
    knn.fit(X_train,y_train)
    y_pred = knn.predict(X_test)
    score.append(accuracy_score(y_test,y_pred))
```

C:\Users\WOLF\anaconda3\lib\site-packages\sklearn\neighbors_classification.p
y:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\WOLF\anaconda3\lib\site-packages\sklearn\neighbors_classification.p
y:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\WOLF\anaconda3\lib\site-packages\sklearn\neighbors_classification.p
y:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along.

In [76]:

```
acc5 = accuracy_score(y_test,y_pred5)
acc5
```

Out[76]:

0.8949044585987261

In [77]:

```
# precision

ps5 = precision_score(y_test,y_pred5)
ps5
```

Out[77]:

0.9090909090909091

In [78]:

```
# recall

rs5 = recall_score(y_test,y_pred5)
rs5
```

Out[78]:

0.8805031446540881

In [79]:

```
# f1_score

f1_s5 = f1_score(y_test,y_pred5)
f1_s5
```

Out[79]:

0.8945686900958466

Gradient Boosting Classifier

In [80]:

```
# creating instance

gbc = GradientBoostingClassifier()

# fit on data

gbc.fit(X_train,y_train)
```

Out[80]:

GradientBoostingClassifier()

In [81]:

```
# prediction

y_pred6 = gbc.predict(X_test)

# accuracy

acc6 = accuracy_score(y_test,y_pred)
acc6
```

Out[81]:

0.8949044585987261

In [82]:

```
# precision

ps6 = precision_score(y_test,y_pred6)
ps6
```

Out[82]:

0.9655172413793104

In [83]:

```
# recall

rs6 = recall_score(y_test,y_pred6)
rs6
```

Out[83]:

0.8805031446540881

In [84]:

f1 score

```
f1_s6 = f1_score(y_test,y_pred6)
f1_s6
```

Out[84]:

0.9210526315789475

In [85]:

```
final_df = pd.DataFrame({'Models': ['LR', 'SVC', 'RF', 'DT', 'KNN', 'GBC'],
                        'ACC': [acc1, acc2, acc3, acc4, acc5, acc6],
                        'PS': [ps1, ps2, ps3, ps4, ps5, ps6],
                        'RS': [rs1, rs2, rs3, rs4, rs5, rs6],
                        'F1': [f1_s1, f1_s2, f1_s3, f1_s4, f1_s5, f1_s5]})
```

In [86]:

final_df

Out[86]:


| | Models | ACC | PS | RS | F1 |
|---|--------|----------|----------|----------|----------|
| 0 | LR | 0.878981 | 0.906040 | 0.849057 | 0.876623 |
| 1 | SVC | 0.894904 | 0.925676 | 0.861635 | 0.892508 |
| 2 | RF | 0.942675 | 0.979592 | 0.905660 | 0.941176 |
| 3 | DT | 0.945860 | 0.955128 | 0.937107 | 0.946032 |
| 4 | KNN | 0.894904 | 0.909091 | 0.880503 | 0.894569 |
| 5 | GBC | 0.894904 | 0.965517 | 0.880503 | 0.894569 |

In [87]:

```
sns.barplot(final_df['Models'],final_df['ACC'])  
plt.title('MODELS VS ACCURACY')  
plt.show()
```

C:\Users\WOLF\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

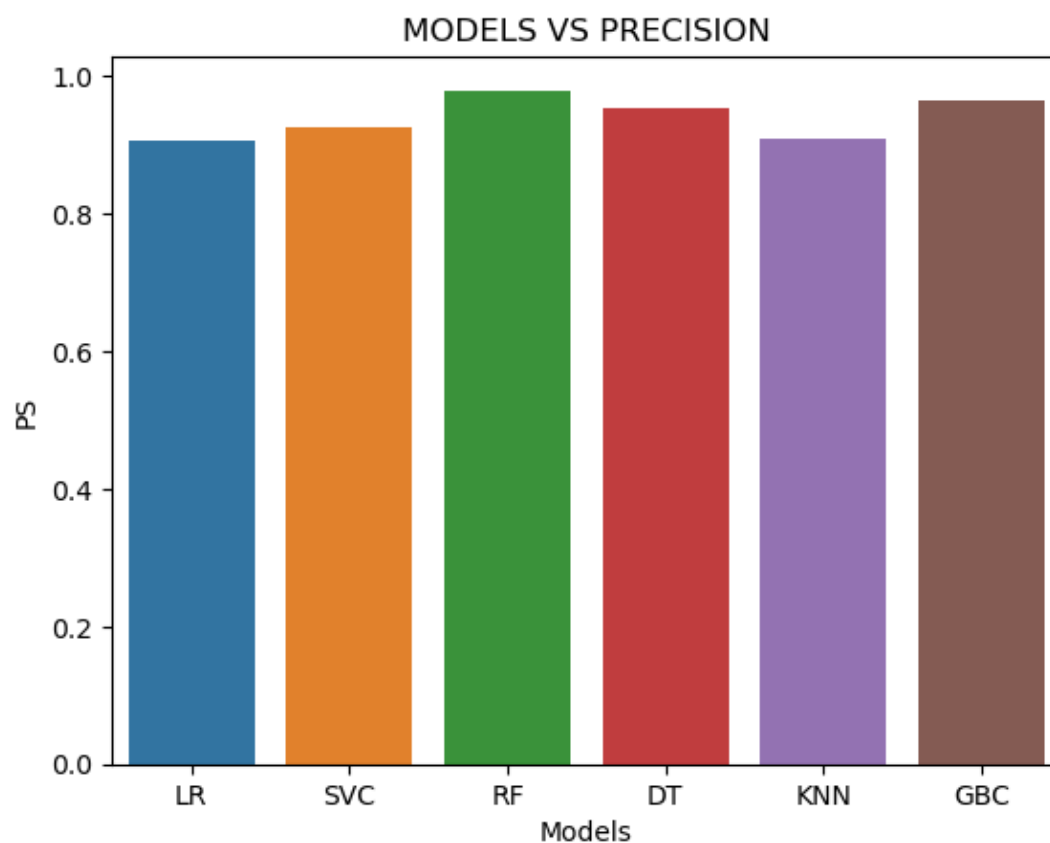


In [88]:

```
sns.barplot(final_df['Models'],final_df['PS'])  
plt.title('MODELS VS PRECISION')  
plt.show()
```

C:\Users\WOLF\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(




In [89]:

```
sns.barplot(final_df['Models'],final_df['RS'])  
plt.title('MODELS VS RECALL')  
plt.show()
```

C:\Users\WOLF\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


warnings.warn(



In [90]:

```
sns.barplot(final_df['Models'],final_df['F1'])  
plt.title('MODELS VS F1 SCORE')  
plt.show()
```

C:\Users\WOLF\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(



As we can see that Decision Tree Classifier is working well for this data set

therefore, we will save this model with DTC

In [91]:

```
# model saving  
  
X_res = std.fit_transform(X_res)
```

In [92]:

```
dtc.fit(X_res,y_res)
```

Out[92]:

```
DecisionTreeClassifier()
```

Saving model using joblib library

In [93]:

```
import joblib
```

In [94]:

```
# first we need to use dump function of library  
# then provide the chosen algorithm and the name we want to save the model
```

```
joblib.dump(dtc, 'insurance_claim_prediction_model')
```

Out[94]:

```
['insurance_claim_prediction_model']
```

In [95]:

```
# to check wether trh emodel is saved or not
```

```
model = joblib.load('insurance_claim_prediction_model')
```

In [96]:

```
# now to verify the model is working or not we will provide it values of first row  
# then cross check wether its working fine or not
```

```
df.head(1)
```

Out[96]:

| | age | sex | bmi | children | smoker | region | charges | insuranceclaim |
|---|-----|-----|------|----------|--------|--------|-----------|----------------|
| 0 | 19 | 0 | 27.9 | 0 | 1 | 3 | 16884.924 | 1 |

In [97]:

```
model.predict([[19,0,27.9,0,1,3,16887.924]])
```

Out[97]:

```
array([1], dtype=int64)
```

insight :

As we can see our model's prediction is perfectly fine

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: